

# Linux File I/O

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Introduce Linux File I/O**

**Buffered I/O vs system calls**

**Opening Files**

# File I/O in Linux

- Files are opened via file descriptors (fd)
  - represented by an integer
- Kernel maintains a Per-process list of files - file table
  - List indexed by fds

# File I/O in Linux

- Every process has at least 3 file descriptors open
  - stdin (fd 0 or STDIN\_FILENO)
    - terminal input device
  - stdout (fd 1 or STDOUT\_FILENO)
    - terminal display
  - stderr (fd 2 or STDERR\_FILENO)
    - terminal display

# Opening and Accessing Files

- `open()` (possibly with `O_CREAT`)
- `read()`
- `write()`
- `close()`
- What about `fopen`, `fwrite`, `fread`?

# Opening and Accessing Files

- What about `fopen`, `fwrite`, `fread`?
  - “f” prefix relate to buffered I/O, not system calls
  - non “f” prefix versions listed here are the underlying system calls.

# fopen vs open()

## FOPEN(3)

Linux Programmer's Manual

### NAME

fopen, fdopen, freopen - stream open functions

### SYNOPSIS

```
#include <stdio.h>
```

```
FILE *fopen(const char *pathname, const char *mode);
```

- Buffered IO uses **FILE\***

## OPEN(2)

Linux Programmer's Manual

### NAME

open, openat, creat - open and possibly create a file

### SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

```
int open(const char *pathname, int flags, mode_t mode);
```

- syscall open() uses file descriptor

# fopen vs open()

## FOPEN(3)

Linux Programmer's Manual

### NAME

fopen, fdopen, freopen - stream open functions

### INTRO(3)

Linux Programmer's Manual

INTRO(3)

### NAME

intro - introduction to library functions

### DESCRIPTION

Section 3 of the manual describes all library functions excluding the library functions (system call wrappers) described in Section 2, which implement system calls.

## OPEN(2)

Linux Programmer's Manual

### NAME

open, openat, creat - open and possibly create a file

### INTRO(2)

Linux Programmer's Manual

### NAME

intro - introduction to system calls

### DESCRIPTION

Section 2 of the manual describes the Linux system calls. A system call is an are not invoked directly: instead, most system calls have corresponding C library trapping to kernel mode) in order to invoke the system call. Thus, making a sys tion.

o Located in different man page sections



# fopen vs open()

- `openat()` is eventually called from `fopen()`

```
extern FILE *_IO_new_fopen (const char*, const char*);
#   define fopen(fname, mode) _IO_new_fopen (fname, mode)
```

```
FILE *
__fopen_internal (const char *filename, const char *mode, int is32)
{
```

```
    _IO_no_init (&new_f->fp.file, 0, 0, &new_f->wd, &_IO_wfile_jumps);
    _IO_JUMPS (&new_f->fp) = &_IO_file_jumps;
    _IO_new_file_init_internal (&new_f->fp);
    if (_IO_file_fopen ((FILE *) new_f, filename, mode, is32) != NULL)
        return __fopen_maybe_mmap (&new_f->fp.file);
```

```
FILE *
_IO_new_fopen (const char *filename, const char *mode)
{
    return __fopen_internal (filename, mode, 1);
}
```

```
FILE *
_IO_file_open (FILE *fp, const char *filename, int posix_mode, int prot,
               int read_write, int is32not64)
{
    int fdesc;
    if (__glibc_unlikely (fp->flags2 & _IO_FLAGS2_NOTCANCEL))
        fdesc = __open_nocancel (filename,
                                posix_mode | (is32not64 ? 0 : O_LARGEFILE), prot);
    else
        fdesc = _open (filename, posix_mode | (is32not64 ? 0 : O_LARGEFILE), prot);
    if (fdesc < 0)
        return NULL;
    fp->file = fdesc;
```

```
return SYSCALL_CANCEL (openat, AT_FDCWD, file, oflag, mode);
```

<https://github.com/bminor/glibc/blob/master/include/stdio.h>  
<https://github.com/bminor/glibc/blob/master/libio/iofopen.c>  
<https://github.com/bminor/glibc/blob/master/libio/fileops.c>  
<https://github.com/bminor/glibc/blob/master/sysdeps/unix/sysv/linux/open.c>

# fopen vs open()

- What's the difference between `fopen()` and `open()`?
  - `f` prefix functions from `<stdio.h>` provide platform-independent user-buffering.
- What is file buffering?
  - Rather than reading/writing a full block, read from or write to copies in memory.
  - Improves performance.

# fopen vs open()

- What do we mean by user-buffering?
  - Happens in userspace, rather than in the kernel.
  - Reduce the amount of system calls into the kernel.



# Opening Files

- `open()`
  - maps pathname to file descriptor
  - Include open flags argument
    - access - (read only, read/write create, append, synchronous I/O, etc)
  - mode argument used when creating
    - read/write/execute permissions for user/group/other

## NAME

`open`, `openat`, `creat` - open and possibly create a file

## SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
```